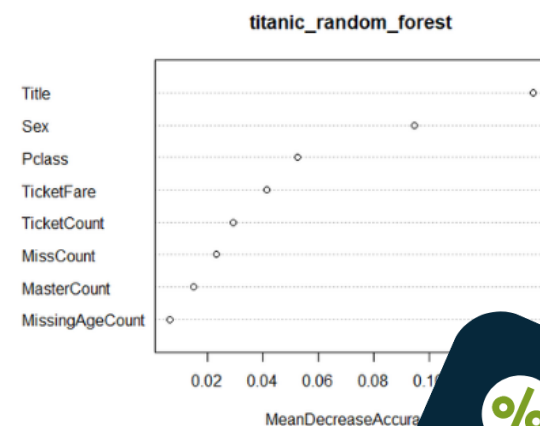
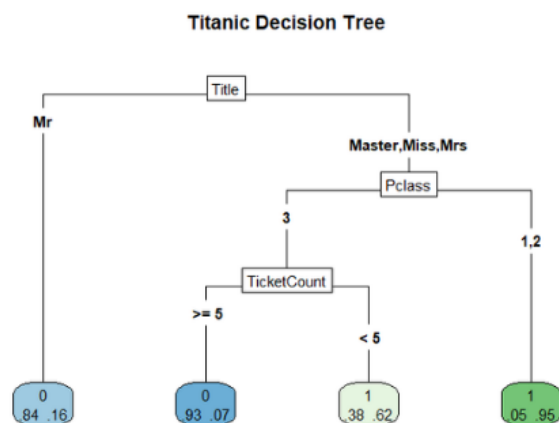



Dave
ON DATA

Introduction to Machine Learning With R

A Course Designed for ANY Professional



**40%
OFF!**

<https://bit.ly/DaveOnDataBlackFriday>

**BLACK
FRIDAY
SALE**

ENDS NOV 29TH



Want to have more impact at work
with data?

Need help applying data analytics
to your business problem?

For students of my online and live courses, I will be offering
personalized 1-on-1 coaching services.

- R programming coaching
- R & Excel data analysis coaching
- Machine learning coaching



CART

Classification Trees

Classification Tree Intuition

Trees Are Rules

Classification trees embody a series of rules to assign/predict *labels*.

Take the following sample of the *Adult Census* data...

Feature/Variables

	occupation	relationship	income
1	Adm-clerical	Not-in-family	<=50K
2	Exec-managerial	Husband	<=50K
3	Handlers-cleaners	Not-in-family	<=50K
4	Handlers-cleaners	Husband	<=50K
5	Prof-specialty	Wife	<=50K
6	Exec-managerial	Wife	<=50K
7	Other-service	Not-in-family	<=50K
8	Exec-managerial	Husband	>50K
9	Prof-specialty	Not-in-family	>50K
10	Exec-managerial	Husband	>50K

Labels

Using the data to the left, consider the following “rules”:

IF occupation **IS IN** (“Adm-clerical”, “Handlers-cleaners”, “Other-service”)
THEN income = “<=50K”

ELSE IF relationship = “Wife” **THEN** income = “<=50K”

ELSE income = “>50K”

How does the classification tree algorithm arrive at these rules?

Classification Trees Minimize “Impurity”

Machine learning algorithms work by trying to achieve an *objective*.

In the case of classification trees, the objective is to *minimize impurity*.

Later we will learn about the math of impurity, right now it’s intuition time...

Feature Values

Label Counts

	occupation	LTE50K	GT50K
1	Adm-clerical	1	0
2	Exec-managerial	2	2
3	Handlers-cleaners	2	0
4	Other-service	1	0
5	Prof-specialty	1	1

These are “pure”

Feature Values

Label Counts

	relationship	LTE50K	GT50K
1	Husband	2	2
2	Not-in-family	3	1
3	Wife	2	0

These are “impure”

The classification tree algorithm iteratively uses the features to split the data into the largest collections of “purest” labels possible.

Let's Build a Tree!

In this contrived example we have 10 *observations* and 2 *features*.

The classification tree algorithm loves a single feature with lots of observations and only a single *label*.

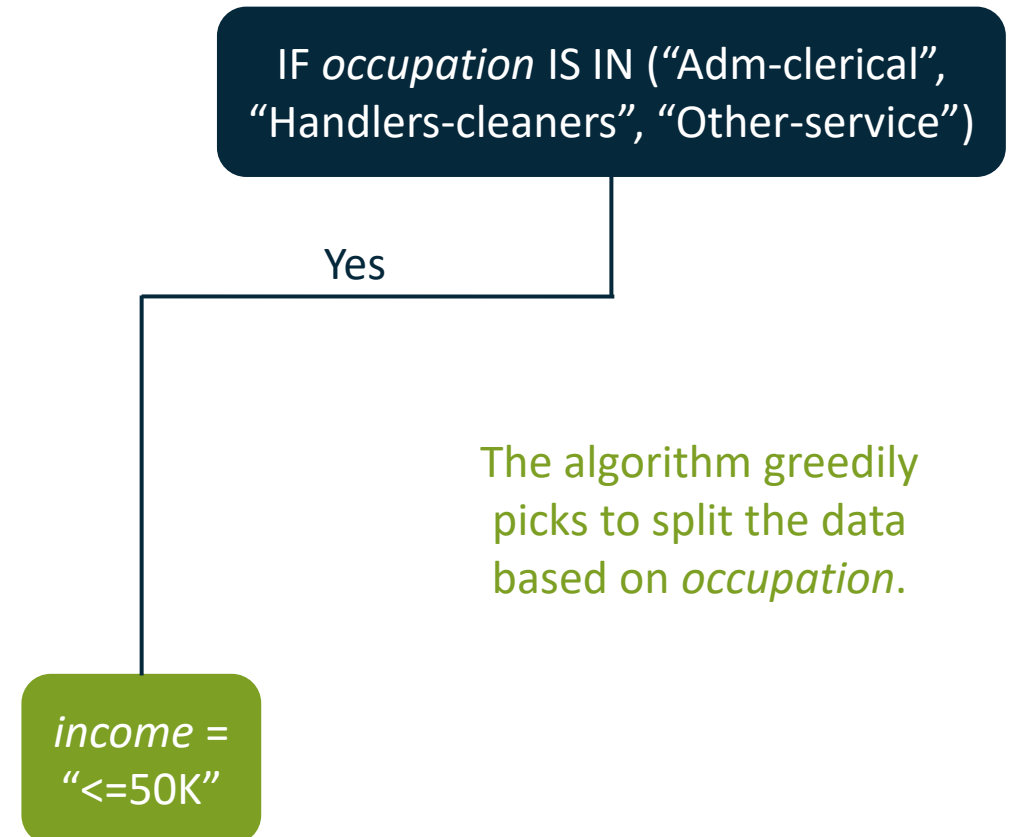
This obsessive love is known as being *greedy*...

	occupation	LTE50K	GT50K
1	Adm-clerical	1	0
2	Exec-managerial	2	2
3	Handlers-cleaners	2	0
4	Other-service	1	0
5	Prof-specialty	1	1

With the *occupation* feature we get 4 observations all with the label “≤50K”

	relationship	LTE50K	GT50K
1	Husband	2	2
2	Not-in-family	3	1
3	Wife	2	0

With the *relationship* feature we get 2 observations all with the label “≤50K”



What's Next?

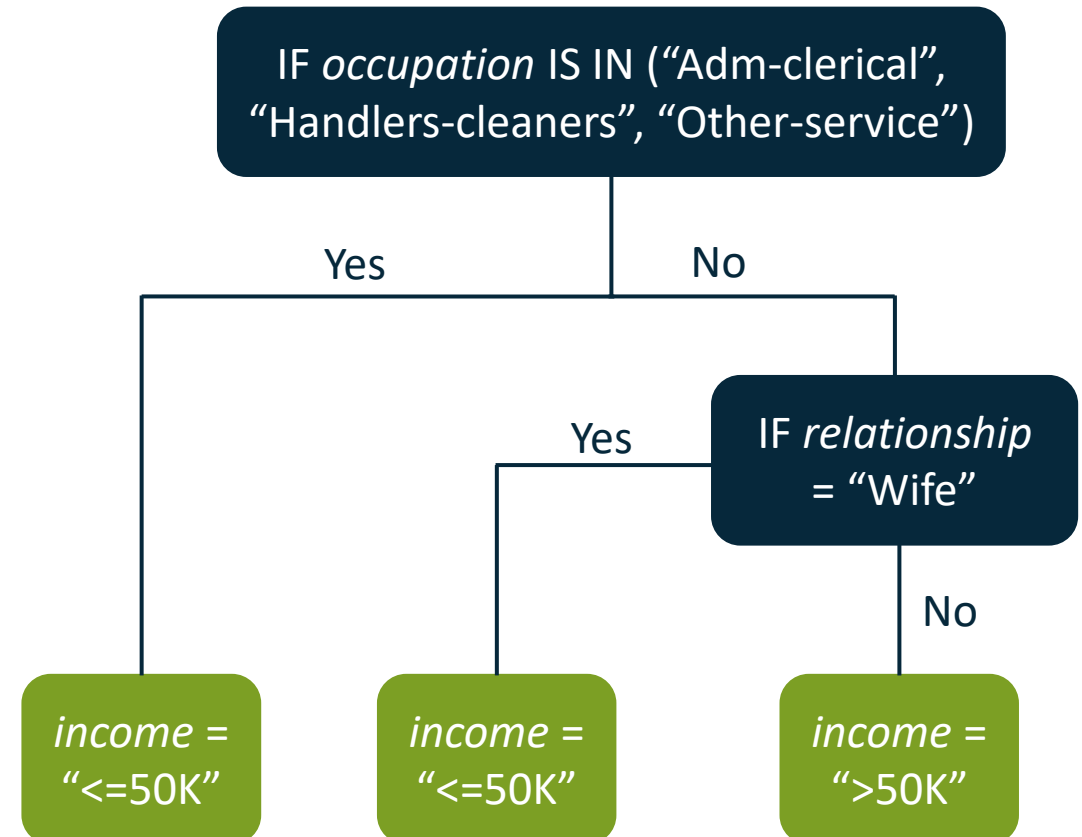
The algorithm has used 40% of the data in the first split...

What's next?

	occupation	relationship	income
1			
2	Exec-managerial	Husband	<=50K
3			
4			
5	Prof-specialty	Wife	<=50K
6	Exec-managerial	Wife	<=50K
7			
8	Exec-managerial	Husband	>50K
9	Prof-specialty	Not-in-family	>50K
10	Exec-managerial	Husband	>50K

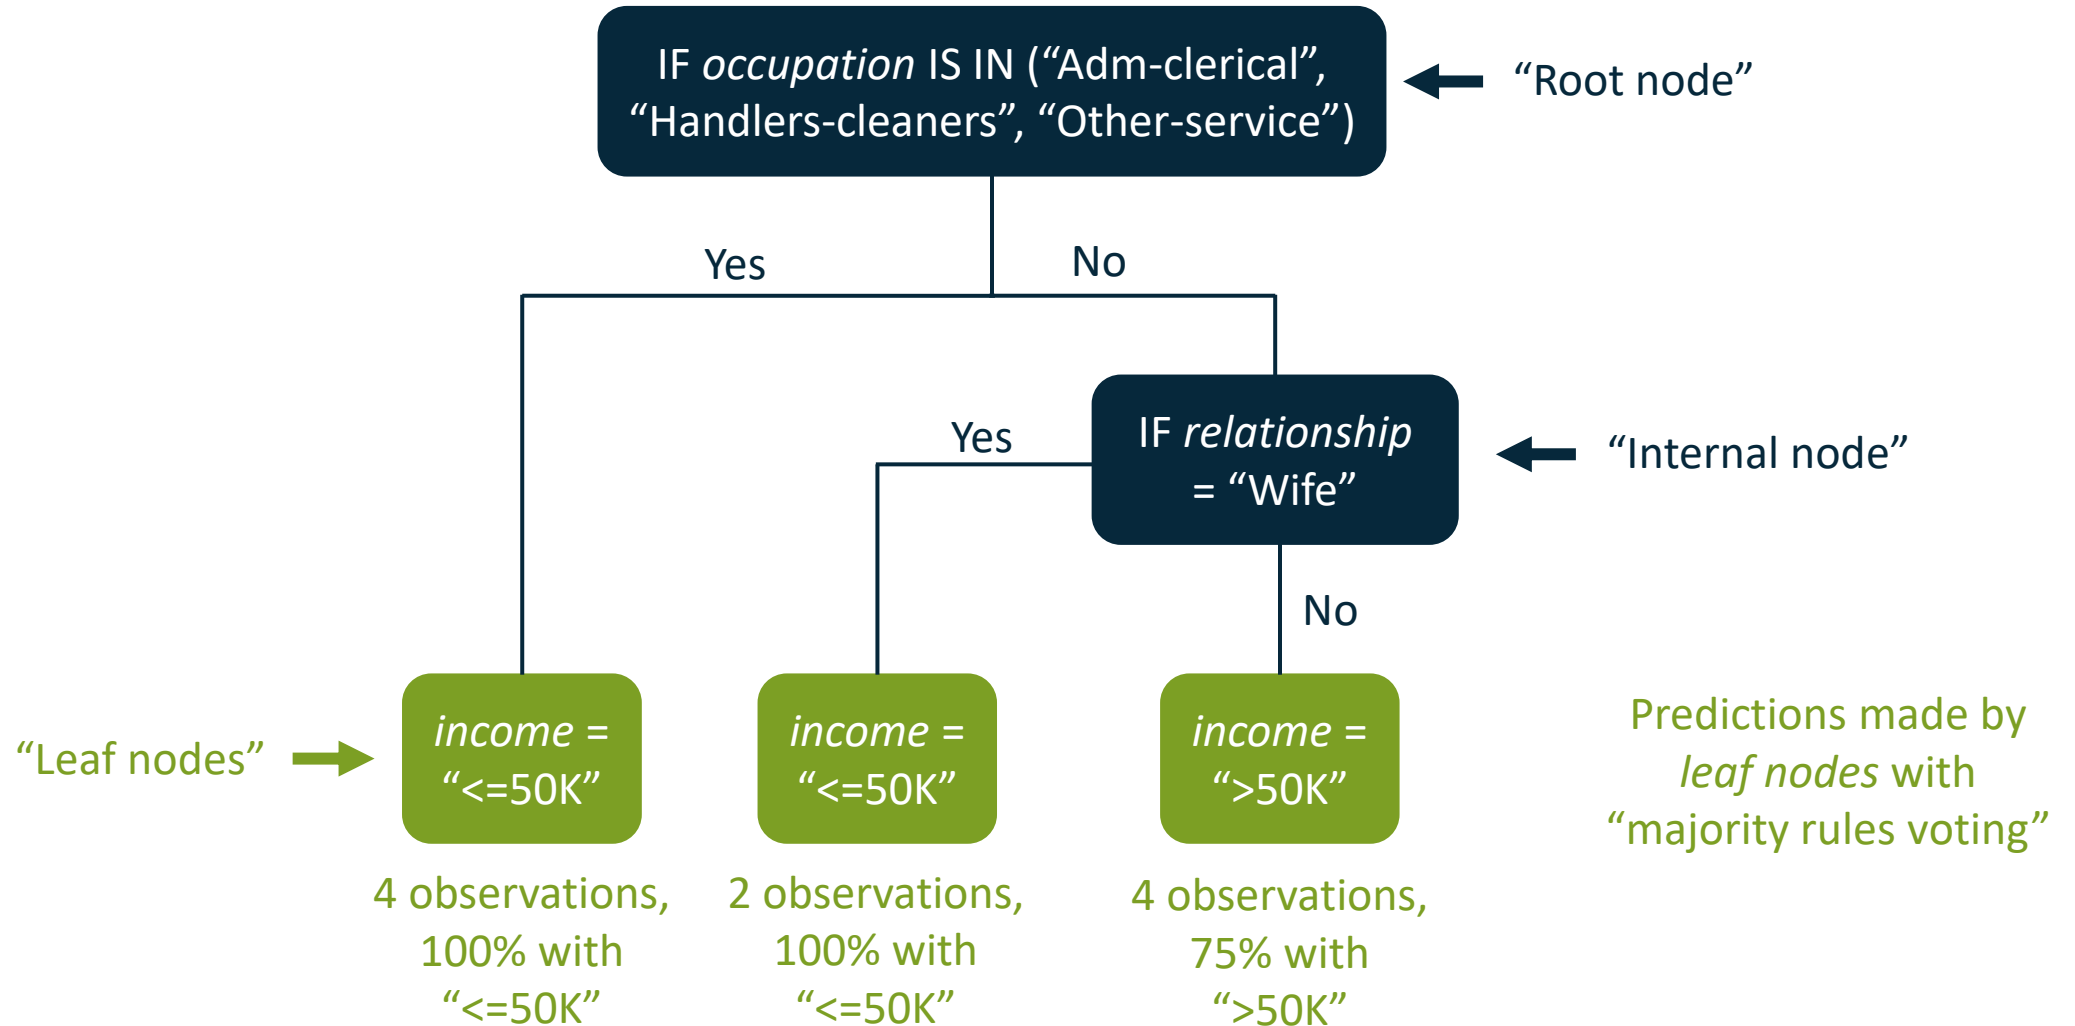
We know this
split is pure!

Of the data remaining, the “impurity math”
we’ll learn later tells us there are no splits left.



Getting Meta

Some things we need to know about trees...



Overfitting Intuition

The Bugbear of Machine Learning

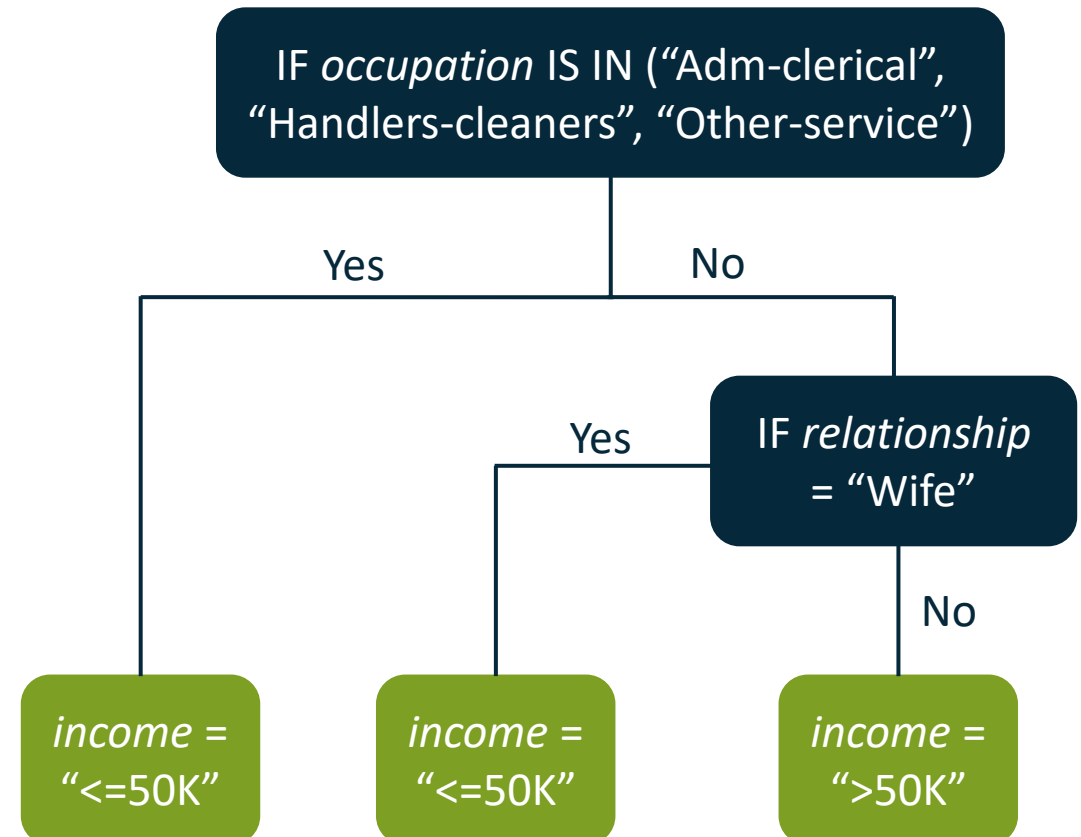
As a professional applying machine learning to your data, you must be paranoid about *overfitting*.

Simply put, overfitting is where your model's predictions are much less "accurate" on new data.

We will be covering overfitting in depth in the next section, for now it's intuition time...

	occupation	relationship	income
1	Adm-clerical	Not-in-family	<=50K
2	Exec-managerial	Husband	<=50K
3	Handlers-cleaners	Not-in-family	<=50K
4	Handlers-cleaners	Husband	<=50K
5	Prof-specialty	Wife	<=50K
6	Exec-managerial	Wife	<=50K
7	Other-service	Not-in-family	<=50K
8	Exec-managerial	Husband	>50K
9	Prof-specialty	Not-in-family	>50K
10	Exec-managerial	Husband	>50K

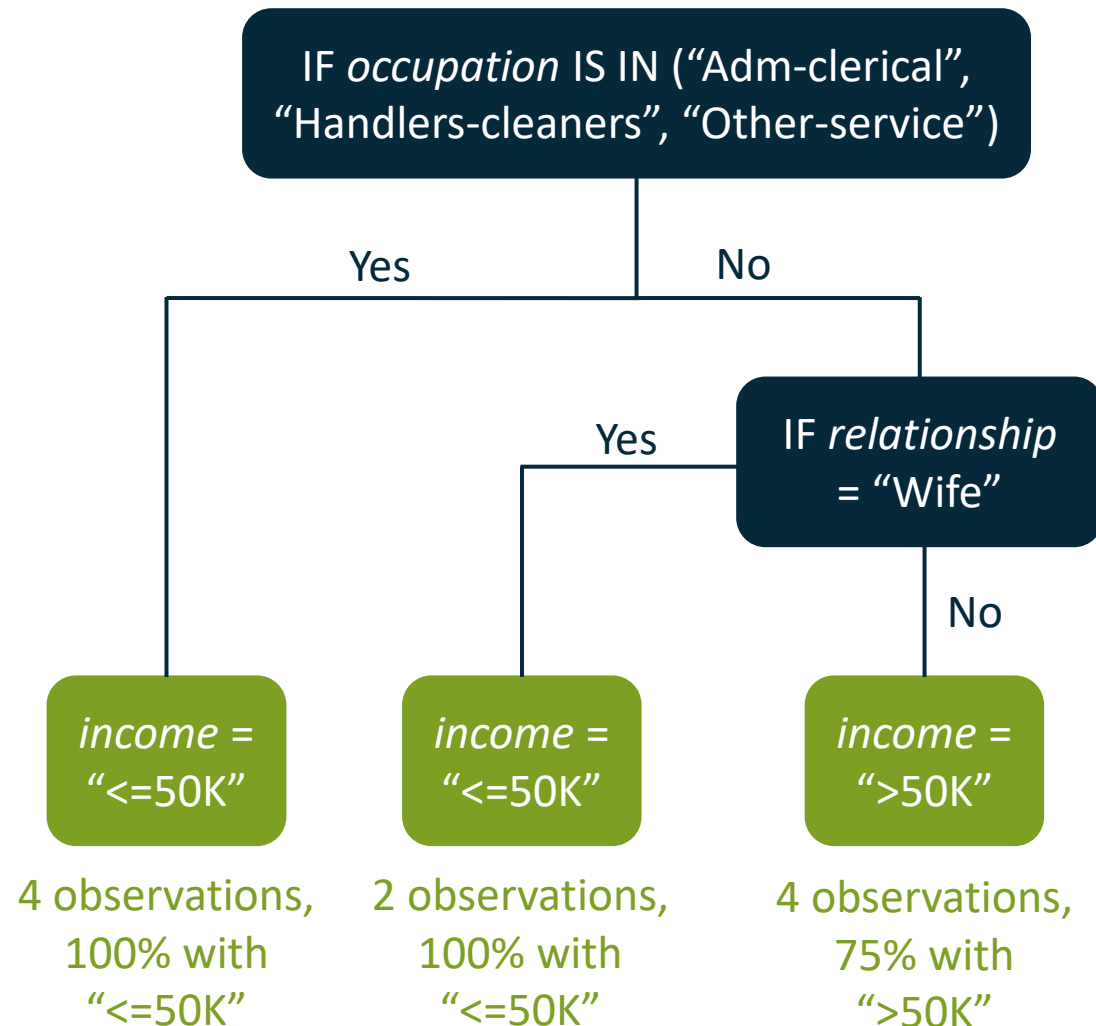
The combination of *data*, *algorithm*, and *training regimen* produced the following model...



The Model Is Good!

Considering the data used to *train* the model, things look awesome!

	occupation	relationship	income
1	Adm-clerical	Not-in-family	<=50K
2	Exec-managerial	Husband	<=50K
3	Handlers-cleaners	Not-in-family	<=50K
4	Handlers-cleaners	Husband	<=50K
5	Prof-specialty	Wife	<=50K
6	Exec-managerial	Wife	<=50K
7	Other-service	Not-in-family	<=50K
8	Exec-managerial	Husband	>50K
9	Prof-specialty	Not-in-family	>50K
10	Exec-managerial	Husband	>50K



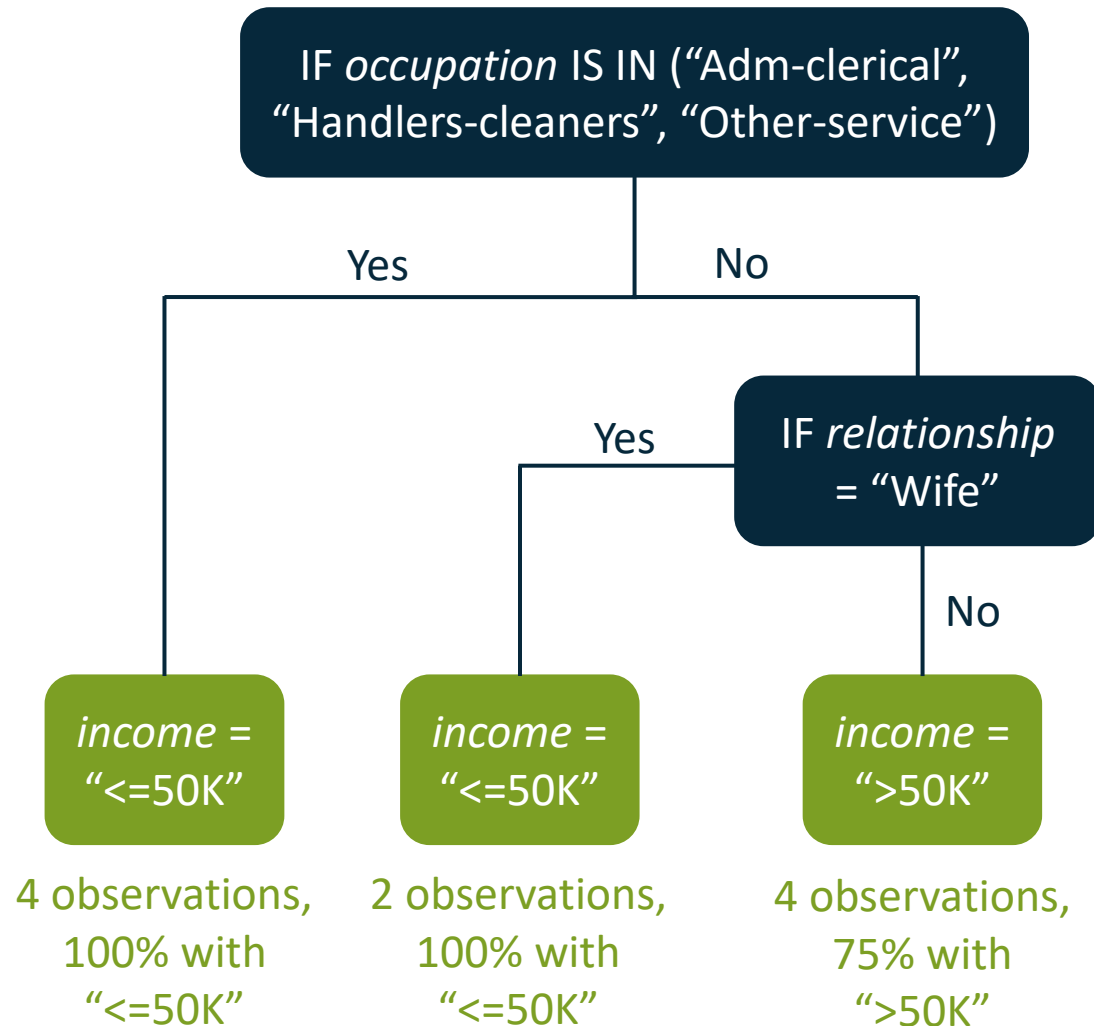
This model
gets 9 out of 10
(i.e., 90%) of
labels correct!

Or Is It?

Consider the following data that wasn't used in training the model...

	occupation	relationship	income
1	Prof-specialty	Wife	>50K
2	Adm-clerical	Wife	>50K
3	Exec-managerial	Wife	>50K
4	Other-service	Husband	>50K
5	Other-service	Husband	>50K
6	Other-service	Wife	>50K
7	Exec-managerial	Husband	<=50K
8	Sales	Not-in-family	<=50K
9	Transport-moving	Husband	<=50K

This is the essence
of overfitting



This model
gets 9 out of 9
(i.e., 100%) of
labels wrong!

What Happened?

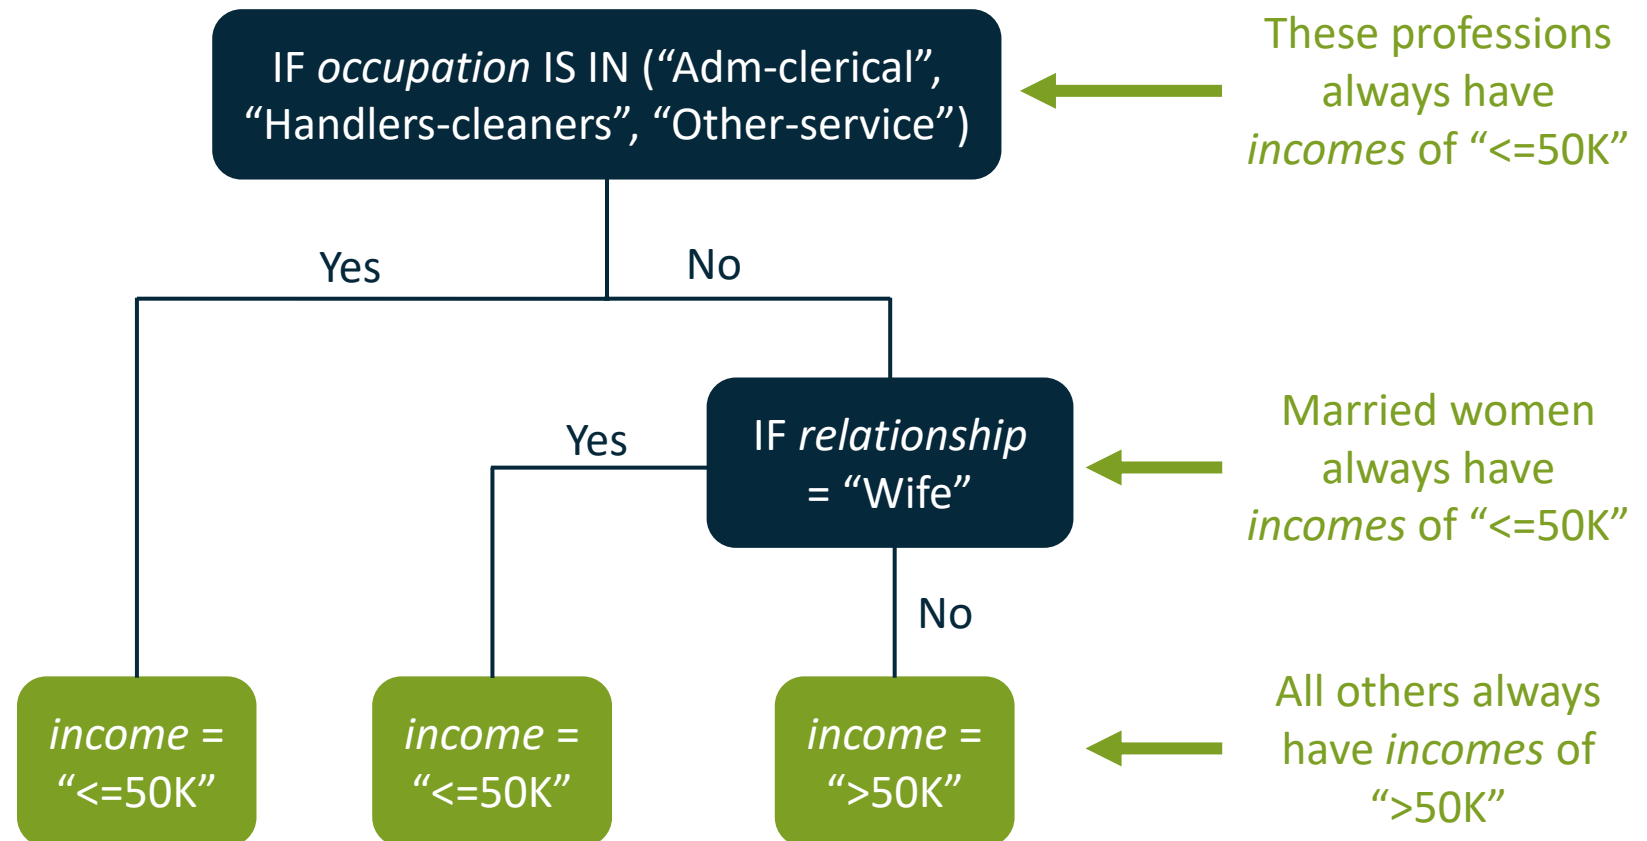
In this example, the overfitting can be blamed on the *training regimen*...

Given this small
amount of data...

	occupation	relationship	income
1	Adm-clerical	Not-in-family	<=50K
2	Exec-managerial	Husband	<=50K
3	Handlers-cleaners	Not-in-family	<=50K
4	Handlers-cleaners	Husband	<=50K
5	Prof-specialty	Wife	<=50K
6	Exec-managerial	Wife	<=50K
7	Other-service	Not-in-family	<=50K
8	Exec-managerial	Husband	>50K
9	Prof-specialty	Not-in-family	>50K
10	Exec-managerial	Husband	>50K

The greediness of
decision trees makes
them prone to overfit!

This model is far too
specialized (or *complex*)!



Model Tuning Intuition

The goal of your *training regimen* is to *tune* your models to combat *overfitting*.

Conceptually, think of it like tuning your car for optimal performance.

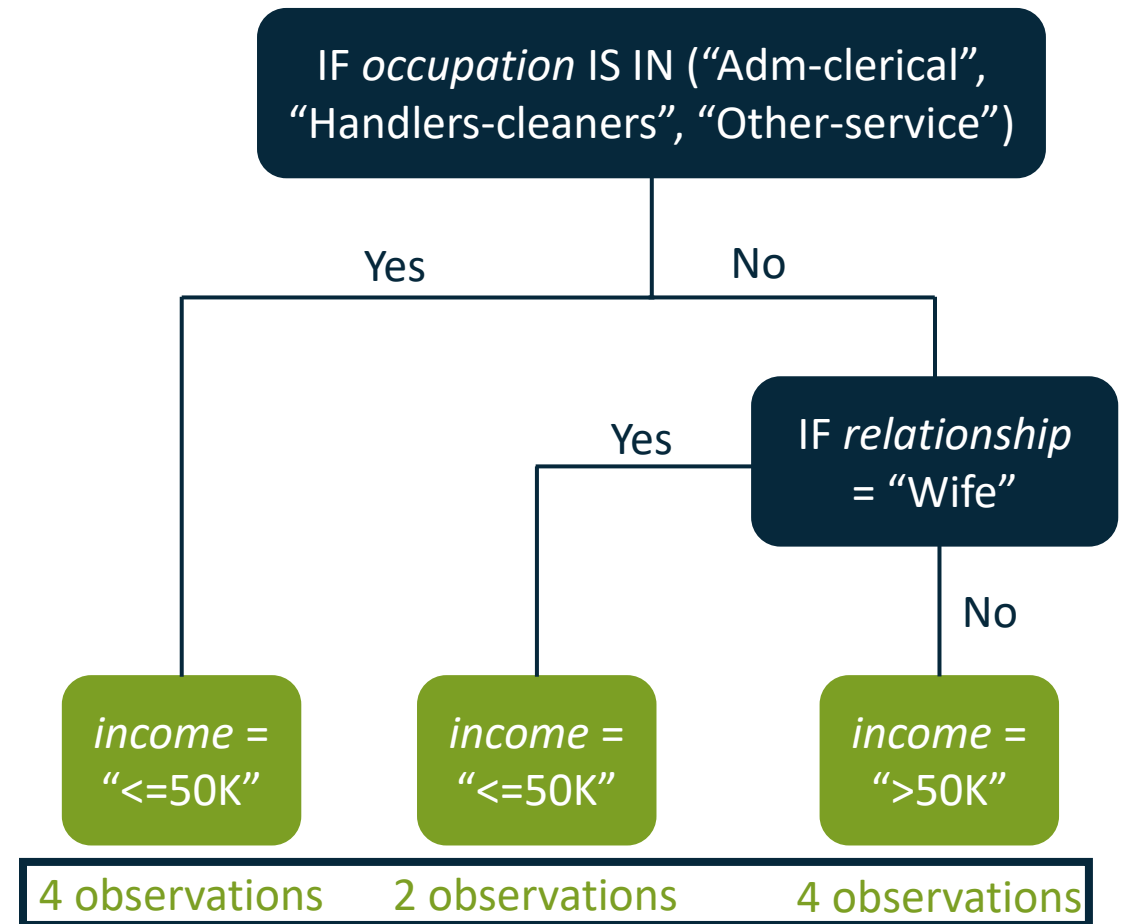
In terms of decisions trees...

One aspect of the decision tree “engine” we can tune is the minimum number of observations required for *leaf nodes*.

Small numbers of leaf node observations allow for more specialized (i.e., *complex*) trees

The minimum number of leaf node observations allowed is known as a *hyperparameter* and can be *tuned* to combat *overfitting*.

We will see later how *algorithms* combat overfitting.



Gini Impurity


Impurity Intuition

We'll start with the 2-label (i.e., the *binary*) scenario of thinking about *impurity*.

Using the *Adult Census* labels as an example, let's think about “buckets” of labels.

We can think of impurity as a spectrum...

The classification tree algorithm needs a calculation that embodies this spectrum to allow it to evaluate each data split.



Gini Impurity

Turns out there are several calculations that manifest the intuition of the previous slide.

In this course we'll use the *Gini impurity* calculation.

Gini impurity is widely used and is the default calculation used in the R packages used in this course.

Don't panic! Here's the equation...

The diagram shows the Gini impurity equation with several green arrows pointing to different parts of the formula to explain its components:


- An arrow points from the text "Go through all the class labels" to the summation symbol \sum .
- An arrow points from the text "Square the proportion" to the term $[p(i|t)]^2$.
- An arrow points from the text "Subtract the stuff to the right from 1" to the minus sign $-$.
- A bracket under the term $[p(i|t)]^2$ is labeled "Proportion of class labels".

$$Gini(t) = 1 - \sum_{i=1}^c [p(i|t)]^2$$

Gini Impurity Example

As we know, the classification tree algorithm's *objective* is to *minimize impurity*.

Now that we've got the math, let's revisit our buckets of labels...



Gini Change

Minimizing Gini Impurity

OK, we know how to calculate *Gini impurity* and it isn't difficult.

We also know that the classification decision tree algorithm's *objective* is to minimize impurity.

The algorithm is going to *greedily* choose splits that produce *the lowest impurity for the most observations*.

Don't panic! Here's the equation...

$$Gini\ Change = Gini(parent\ node) - \sum_{j=1}^2 \frac{N(v_j)}{N} Gini(v_j)$$

Diagram illustrating the equation for Gini Change:

- "Go through both splits" points to the summation symbol $\sum_{j=1}^2$.
- "Count of observations for the split" points to $N(v_j)$.
- "Count of observations in parent node" points to N .
- "Gini impurity for split" points to $Gini(v_j)$.

Training Our First Classification Tree

We now know what is needed to learn (or *train*) our first classification tree from data...

	college	union	manager	income
1	no	yes	no	>50K
2	no	yes	no	>50K
3	no	no	no	<=50K
4	no	no	no	<=50K
5	no	no	no	<=50K
6	yes	no	yes	>50K
7	yes	no	yes	>50K
8	yes	no	yes	>50K
9	yes	yes	no	<=50K
10	yes	yes	no	<=50K

Step #1 – What’s our beginning Gini impurity?

$$\text{Gini Start} = 1 - (5/10)^2 - (5/10)^2$$

$$\text{Gini Start} = 1 - (0.25) - (0.25)$$

$$\text{Gini Start} = 0.5$$

Step #2 – What’s Gini impurity for *college* based on splits?

$$\text{Gini “yes”} = 1 - (2/5)^2 - (3/5)^2$$

$$\text{Gini “yes”} = 1 - (0.16) - (0.36)$$

$$\text{Gini “yes”} = 0.48$$

$$\text{Gini “no”} = 1 - (3/5)^2 - (2/5)^2$$

$$\text{Gini “no”} = 1 - (0.36) - (0.16)$$

$$\text{Gini “no”} = 0.48$$

Step #3 – Weight the *college* splits by observation counts:

$$\text{Weighted Gini “yes”} = (5/10)(0.48) = 0.24$$

$$\text{Weighted Gini “no”} = (5/10)(0.48) = 0.24$$

Step #4 – Gini change with *college*:

$$\text{Gini Change} = (0.5) - (0.24) - (0.24) = 0.02$$

Training Our First Classification Tree

	college	union	manager	income
1	no	yes	no	>50K
2	no	yes	no	>50K
3	no	no	no	<=50K
4	no	no	no	<=50K
5	no	no	no	<=50K
6	yes	no	yes	>50K
7	yes	no	yes	>50K
8	yes	no	yes	>50K
9	yes	yes	no	<=50K
10	yes	yes	no	<=50K

Step #5 – What's Gini impurity for *union* based on splits?

$$\text{Gini "yes"} = 1 - (2/4)^2 - (2/4)^2$$

$$\text{Gini "yes"} = 1 - (0.25) - (0.25)$$

$$\text{Gini "yes"} = 0.5$$

$$\text{Gini "no"} = 1 - (3/6)^2 - (3/6)^2$$

$$\text{Gini "no"} = 1 - (0.25) - (0.25)$$

$$\text{Gini "no"} = 0.5$$

Step #6 – Weight the *union* splits by observation counts:

$$\text{Weighted Gini "yes"} = (4/10)(0.5) = 0.2$$

$$\text{Weighted Gini "no"} = (6/10)(0.5) = 0.3$$

Step #7 – Gini change with *union*:

$$\text{Gini Change} = (0.5) - (0.2) - (0.3) = 0.0$$

Using the *union* feature to split the data doesn't help things!

Training Our First Classification Tree

	college	union	manager	income
1	no	yes	no	>50K
2	no	yes	no	>50K
3	no	no	no	<=50K
4	no	no	no	<=50K
5	no	no	no	<=50K
6	yes	no	yes	>50K
7	yes	no	yes	>50K
8	yes	no	yes	>50K
9	yes	yes	no	<=50K
10	yes	yes	no	<=50K

Step #8 – What's Gini impurity for *manager* based on splits?

$$\text{Gini "yes"} = 1 - (0/3)^2 - (3/3)^2$$

$$\text{Gini "yes"} = 1 - (0.0) - (1)$$

$$\text{Gini "yes"} = 0.0$$

$$\text{Gini "no"} = 1 - (5/7)^2 - (2/7)^2$$

$$\text{Gini "no"} = 1 - (0.51) - (0.08)$$

$$\text{Gini "no"} = 0.41$$

Step #9 – Weight the *manager* splits by observation counts:

$$\text{Weighted Gini "yes"} = (3/10)(0.0) = 0.0$$

$$\text{Weighted Gini "no"} = (7/10)(0.41) = 0.29$$

Step #10 – Gini change with *manager*:

$$\text{Gini Change} = (0.5) - (0.0) - (0.29) = 0.21$$

Using the *manager* feature to split the data helps the most!

The Tree So Far

Not a *manager*?
manager == "no"

Yes

No

Got some
work to do...

	college	union	manager	income
1	no	yes	no	>50K
2	no	yes	no	>50K
3	no	no	no	<=50K
4	no	no	no	<=50K
5	no	no	no	<=50K
6	yes	yes	no	<=50K
7	yes	yes	no	<=50K

Awesome!

	college	union	manager	income
1	yes	no	yes	>50K
2	yes	no	yes	>50K
3	yes	no	yes	>50K

income =
">50K"

The *manager*
feature is no
longer useful

Training Our First Classification Tree

	college	union	manager	income
1	no	yes	no	>50K
2	no	yes	no	>50K
3	no	no	no	<=50K
4	no	no	no	<=50K
5	no	no	no	<=50K
6	yes	yes	no	<=50K
7	yes	yes	no	<=50K

Step #11 – Beginning Gini impurity:

$$\text{Gini Start} = 1 - (5/7)^2 - (2/7)^2$$

$$\text{Gini Start} = 1 - (0.51) - (0.08) = 0.41$$

Step #12 – What's Gini impurity for *college* based on splits?

$$\text{Gini "yes"} = 1 - (2/2)^2 - (0/2)^2$$

$$\text{Gini "yes"} = 1 - (1) - (0.0)$$

$$\text{Gini "yes"} = 0.0$$

$$\text{Gini "no"} = 1 - (3/5)^2 - (2/5)^2$$

$$\text{Gini "no"} = 1 - (0.36) - (0.16)$$

$$\text{Gini "no"} = 0.48$$

Step #13 – Weight the *college* splits by observation counts:

$$\text{Weighted Gini "yes"} = (2/7)(0.0) = 0.0$$

$$\text{Weighted Gini "no"} = (5/7)(0.48) = 0.34$$

Step #14 – Gini change with *college*:

$$\text{Gini Change} = (0.41) - (0.0) - (0.34) = 0.07$$

Training Our First Classification Tree

	college	union	manager	income
1	no	yes	no	>50K
2	no	yes	no	>50K
3	no	no	no	<=50K
4	no	no	no	<=50K
5	no	no	no	<=50K
6	yes	yes	no	<=50K
7	yes	yes	no	<=50K

Step #15 – What's Gini impurity for *union* based on splits?

$$\text{Gini "yes"} = 1 - (2/4)^2 - (2/4)^2$$

$$\text{Gini "yes"} = 1 - (0.25) - (0.25)$$

$$\text{Gini "yes"} = 0.5$$

$$\text{Gini "no"} = 1 - (3/3)^2 - (0/0)^2$$

$$\text{Gini "no"} = 1 - (1) - (0)$$

$$\text{Gini "no"} = 0.0$$

Step #16 – Weight the *union* splits by observation counts:

$$\text{Weighted Gini "yes"} = (4/7)(0.5) = 0.29$$

$$\text{Weighted Gini "no"} = (3/7)(0.0) = 0.0$$

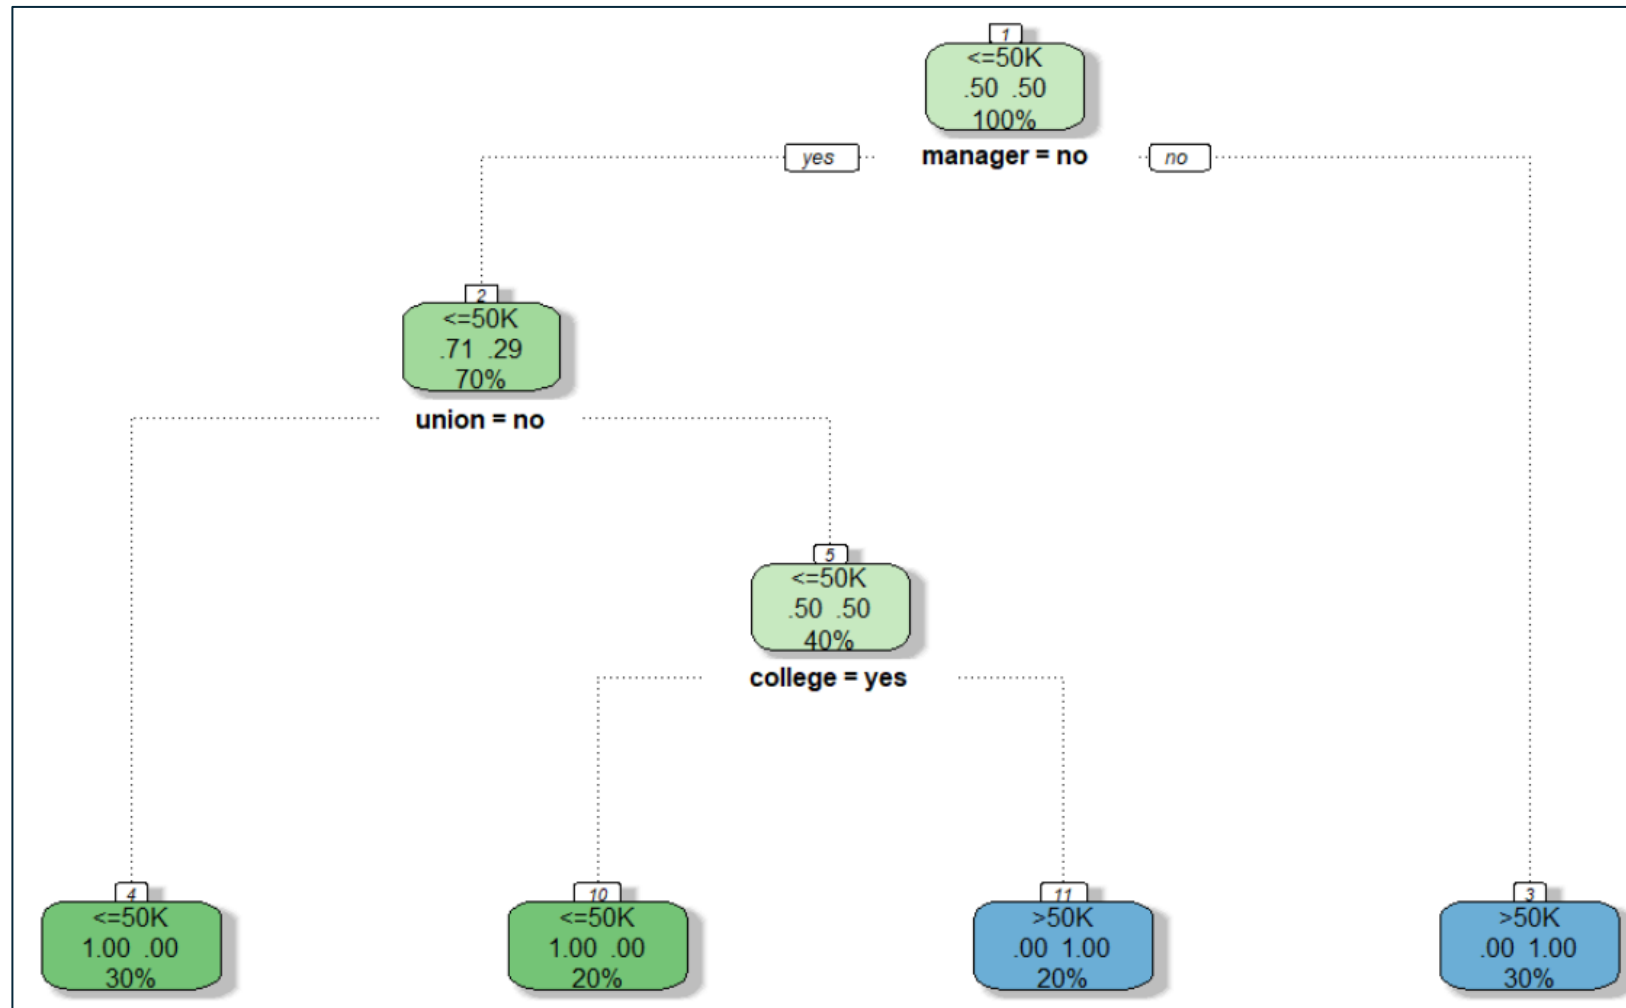
Step #17 – Gini change with *union*:

$$\text{Gini Change} = (0.41) - (0.29) - (0.0) = 0.12$$

Using the *union* feature to split the data helps the most!

The Real Deal

A classification tree trained in R using the contrived data of this example...



Many Categories Impurity

So Many Possibilities

As we know, the CART algorithm uses 2-way (i.e., *binary*) data splits.

So how does this map to a classification tree where a feature has many categories?

For example...

Pivoting the
occupation feature...

	occupation	relationship	income
1	Adm-clerical	Not-in-family	<=50K
2	Exec-managerial	Husband	<=50K
3	Handlers-cleaners	Not-in-family	<=50K
4	Handlers-cleaners	Husband	<=50K
5	Prof-specialty	Wife	<=50K
6	Exec-managerial	Wife	<=50K
7	Other-service	Not-in-family	<=50K
8	Exec-managerial	Husband	>50K
9	Prof-specialty	Not-in-family	>50K
10	Exec-managerial	Husband	>50K

5 categories...

	occupation	category_count
1	Adm-clerical	1
2	Other-service	1
3	Handlers-cleaners	2
4	Prof-specialty	2
5	Exec-managerial	4

So many possibilities!

IF *occupation* IS IN ("Adm-clerical")

IF *occupation* IS IN ("Adm-clerical",
"Other-service")

IF *occupation* IS IN ("Adm-clerical",
"Other-service", "Handlers-cleaners")

IF *occupation* IS IN ("Adm-clerical", "Other-
service", "Handlers-cleaners", "Prof-specialty")

CART Is Smart

As the number of categories increases, the number of possible binary splits explodes.

Luckily, CART is smart and uses an optimization...

Pivoting the
occupation feature...

	occupation	relationship	income
1	Adm-clerical	Not-in-family	<=50K
2	Exec-managerial	Husband	<=50K
3	Handlers-cleaners	Not-in-family	<=50K
4	Handlers-cleaners	Husband	<=50K
5	Prof-specialty	Wife	<=50K
6	Exec-managerial	Wife	<=50K
7	Other-service	Not-in-family	<=50K
8	Exec-managerial	Husband	>50K
9	Prof-specialty	Not-in-family	>50K
10	Exec-managerial	Husband	>50K

Sort the categories by the chance of
income being "<=50K"...

	occupation	Count	LTE50K	GT50K	ChanceLTE
1	Adm-clerical	1	1	0	1.0
2	Handlers-cleaners	2	2	0	1.0
3	Other-service	1	1	0	1.0
4	Exec-managerial	4	2	2	0.5
5	Prof-specialty	2	1	1	0.5

Calculate the Gini change by
adding each category in turn...

IF *occupation* IS IN
("Adm-clerical")

IF *occupation* IS IN
("Adm-clerical",
"Handlers-cleaners")

IF *occupation* IS IN
("Adm-clerical",
"Handlers-cleaners",
"Other-service")

CART only needs to calculate the Gini change
for a tiny subset of possible splits!

An Example

As an example, take the first split in a classification tree *trained* from the following data...

	occupation	relationship	income
1	Adm-clerical	Not-in-family	<=50K
2	Exec-managerial	Husband	<=50K
3	Handlers-cleaners	Not-in-family	<=50K
4	Handlers-cleaners	Husband	<=50K
5	Prof-specialty	Wife	<=50K
6	Exec-managerial	Wife	<=50K
7	Other-service	Not-in-family	<=50K
8	Exec-managerial	Husband	>50K
9	Prof-specialty	Not-in-family	>50K
10	Exec-managerial	Husband	>50K

	occupation	Count	LTE50K	GT50K	ChanceLTE
1	Adm-clerical	1	1	0	1.0
2	Handlers-cleaners	2	2	0	1.0
3	Other-service	1	1	0	1.0
4	Exec-managerial	4	2	2	0.5
5	Prof-specialty	2	1	1	0.5

	relationship	Count	LTE50K	GT50K	ChanceLTE
1	Wife	2	2	0	1.00
2	Not-in-family	4	3	1	0.75
3	Husband	4	2	2	0.50

Gini Start = 0.42

Gini Change = 0.02

Gini Change = 0.08

Gini Change = 0.12

Gini Change = 0.06

Gini Change = 0.05

Gini Change = 0.05

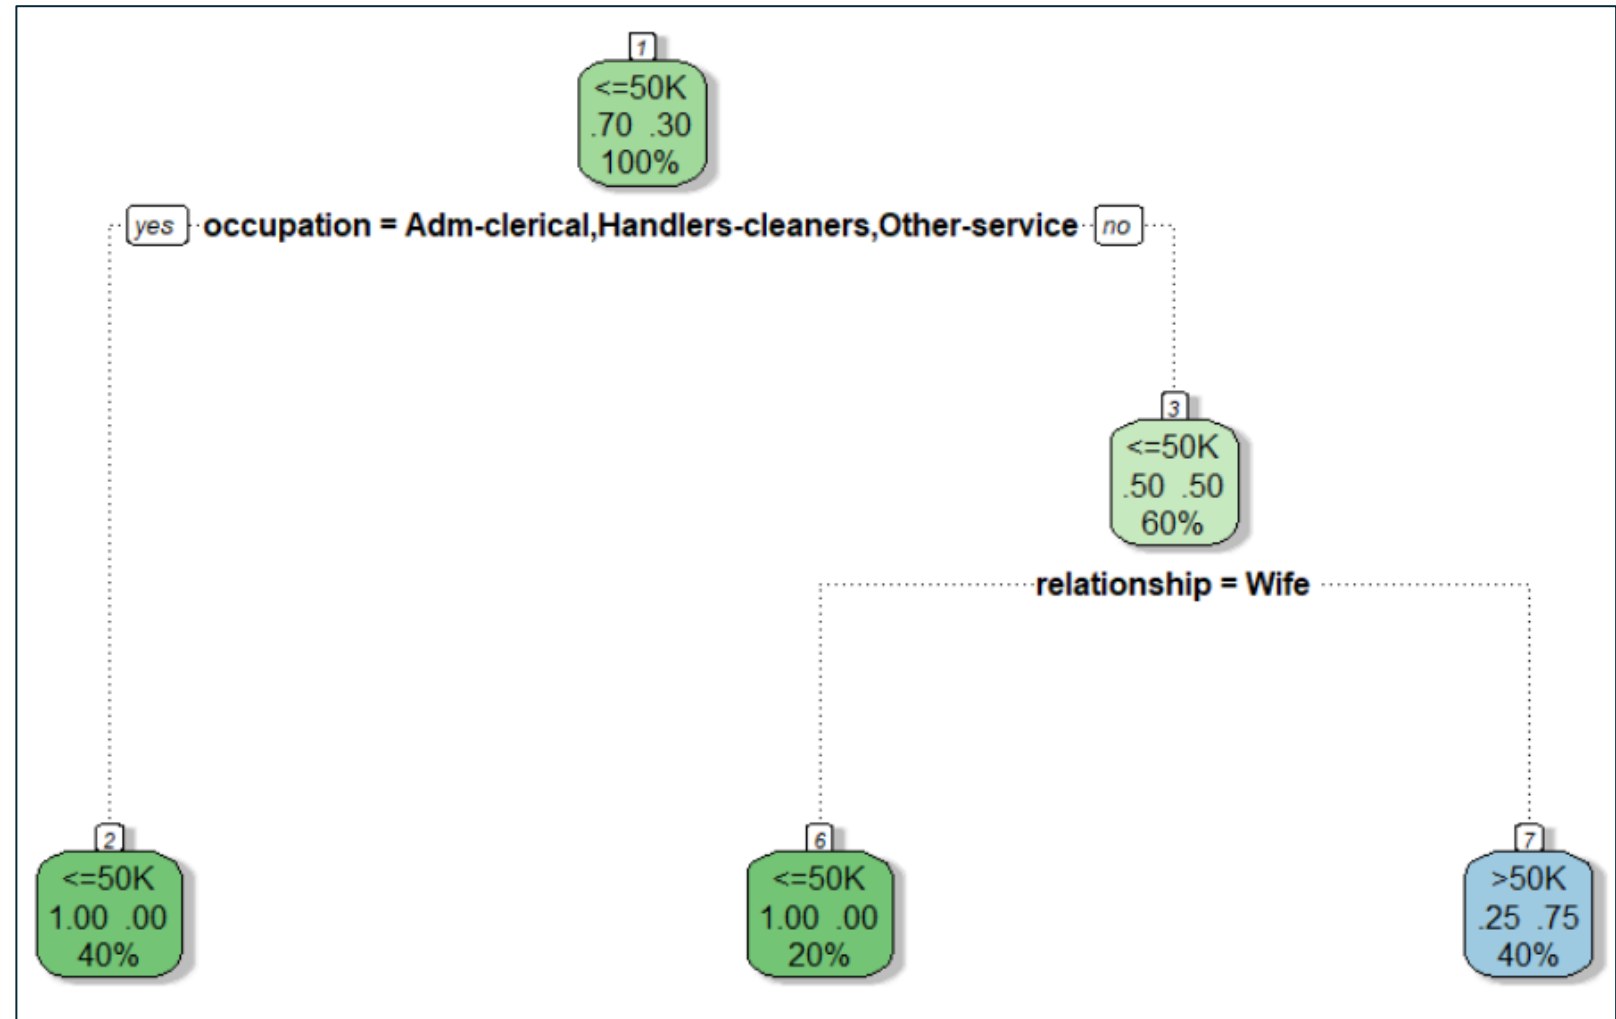
The best first split for the tree:

IF *occupation* IS IN ("Adm-clerical",
"Handlers-cleaners", "Other-service")

The Real Deal

Check out a classification tree trained in R...

	occupation	relationship	income
1	Adm-clerical	Not-in-family	<=50K
2	Exec-managerial	Husband	<=50K
3	Handlers-cleaners	Not-in-family	<=50K
4	Handlers-cleaners	Husband	<=50K
5	Prof-specialty	Wife	<=50K
6	Exec-managerial	Wife	<=50K
7	Other-service	Not-in-family	<=50K
8	Exec-managerial	Husband	>50K
9	Prof-specialty	Not-in-family	>50K
10	Exec-managerial	Husband	>50K



Numeric Feature Impurity

Again, So Many Possibilities

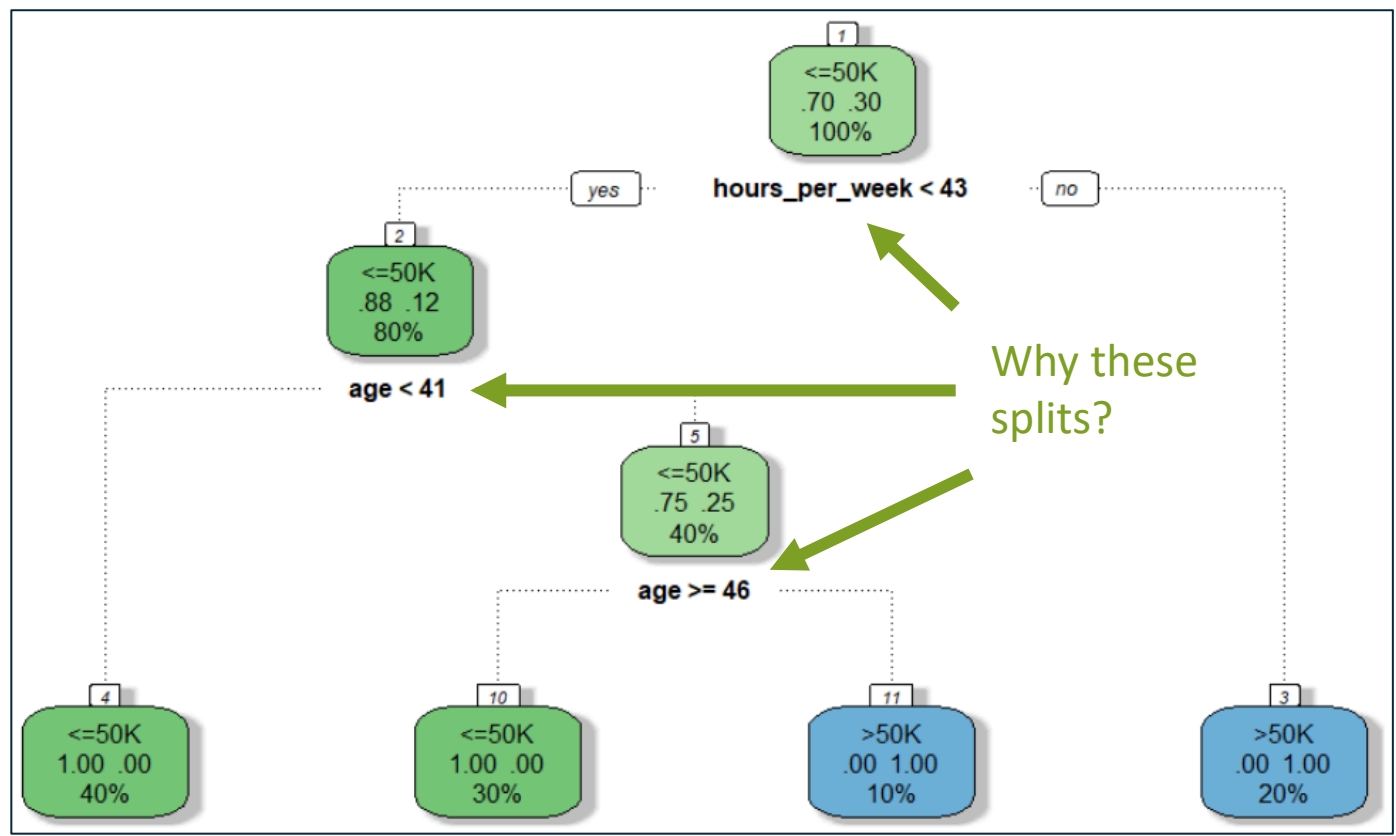
We learned in the last lesson how the CART algorithm finds the best split for features with many categories.

Think about numeric features. Numeric features often have many unique values.

Take the following example from the *Adult Census* data...

Using this data, R trains the following tree...

	hours_per_week	age	income
1	40	39	<=50K
2	13	50	<=50K
3	40	38	<=50K
4	40	53	<=50K
5	40	28	<=50K
6	40	37	<=50K
7	16	49	<=50K
8	45	52	>50K
9	50	31	>50K
10	40	42	>50K



Again, CART Is Smart

In real data sets, numeric features can have many unique values.

Once again, CART is smart and uses an optimization to find optimal splits...

Sort the *hours_per_week* feature...

	hours_per_week	age	income
1	40	39	<=50K
2	13	50	<=50K
3	40	38	<=50K
4	40	53	<=50K
5	40	28	<=50K
6	40	37	<=50K
7	16	49	<=50K
8	45	52	>50K
9	50	31	>50K
10	40	42	>50K

Find the observations where the label changes...

	hours_per_week	age	income
1	13	50	<=50K
2	16	49	<=50K
3	40	39	<=50K
4	40	38	<=50K
5	40	53	<=50K
6	40	28	<=50K
7	40	37	<=50K
8	40	42	>50K
9	45	52	>50K
10	50	31	>50K

Move split point down to where values change...

	hours_per_week	age	income
1	13	50	<=50K
2	16	49	<=50K
3	40	39	<=50K
4	40	38	<=50K
5	40	53	<=50K
6	40	28	<=50K
7	40	37	<=50K
8	40	42	>50K
9	45	52	>50K
10	50	31	>50K

Calculate Gini change by
splitting between the two values

Doh! Values are
the same!

Split = $(40 + 45) / 2 = 42.5 = 43$
Gini Change = 0.245

Continuing...

Sort the *age* feature...

Possible
splits...

	hours_per_week	age	income
1	40	28	<=50K
2	50	31	>50K
3	40	37	<=50K
4	40	38	<=50K
5	40	39	<=50K
6	40	42	>50K
7	16	49	<=50K
8	13	50	<=50K
9	45	52	>50K
10	40	53	<=50K

Gini Change = 0.02
Gini Change = 0.02

Gini Change = 0.02
Gini Change = 0.003

Gini Change = 0.02
Gini Change = 0.02

Process continues with left side
of tree using the *age* feature.

The *hours_per_week* split is optimal...

IF *hours_per_week* < 43

Yes

No

	hours_per_week	age	income
1	40	39	<=50K
2	13	50	<=50K
3	40	38	<=50K
4	40	53	<=50K
5	40	28	<=50K
6	40	37	<=50K
7	16	49	<=50K
8	40	42	>50K

	hours_per_week	age	income
1	45	52	>50K
2	50	31	>50K

income =
">50K"